

# Data 612 - Project 2

Paul Perez

6/16/2020

## Contents

The Recommender System's . . . . .	1
Data Ingestion, Selection, Manipulation . . . . .	2
Data Exploration . . . . .	2
Split into Training & Test datasets using an 80/20 Ratio . . . . .	8
Item-Item Collaborative Filtering . . . . .	8
User-User Collaborative Filtering . . . . .	8
Summary . . . . .	9

## The Recommender System's

This recommender system uses the `recommenderlab` package, and we'll load the MovieLense dataset. We can see that using the `class()` function, that this MovieLense dataset is a `realRatingMatrix`.

## Data Ingestion, Selection, Manipulation

```
data("MovieLense")
class(MovieLense)
```

```
## [1] "realRatingMatrix"
## attr(,"package")
## [1] "recommenderlab"
```

```
matrix <- as(MovieLense, "realRatingMatrix")
```

## Data Exploration

Using the `dim()` function, we can see the dimensions of the MovieLense matrix.

```
dim(matrix)
```


```
## [1] 943 1664
```

There are 943 users and 1664 movies.


## Similarity Matrix (User & Item)

The similarity matrix measures the similarity between item's or user's within the matrix. For the purpose of this project, we'll use the `cosine` method.

```
user_similarity <- similarity(matrix[1:10, ], method = "cosine", which = "users")
image(as.matrix(user_similarity), main = "User Similarity Matrix")
```



```
item_similarity <- similarity(matrix[, 1:10], method = "cosine", which = "items")  
image(as.matrix(item_similarity), main = "Item Similarity Matrix")
```



The darker the cell in each of these similarity matrices show's the greater similarity between the two user's or item's.

There are some emperical studies that show greater performance with the **cosine** similarity method for item based recommendation systems where the **pearson** similarity method shows greater performance for user based recommendation systems.

Since we loaded the data as a **realRatingMatrix**, which is an S4 class, we can look further into components of the object using `slotNames()`

```
slotNames(matrix)
```

```
## [1] "data"      "normalize"
```

By adding `@data` to the `matrix` object, we'll be able to take a look at this class, which happens to be a `dgCMatrix` class that inherits from the original matrix.

```
class(matrix@data)
```

```
## [1] "dgCMatrix"
## attr("package")
## [1] "Matrix"
```

By converting the ratings to a vector, we can count the number of unique ratings available.

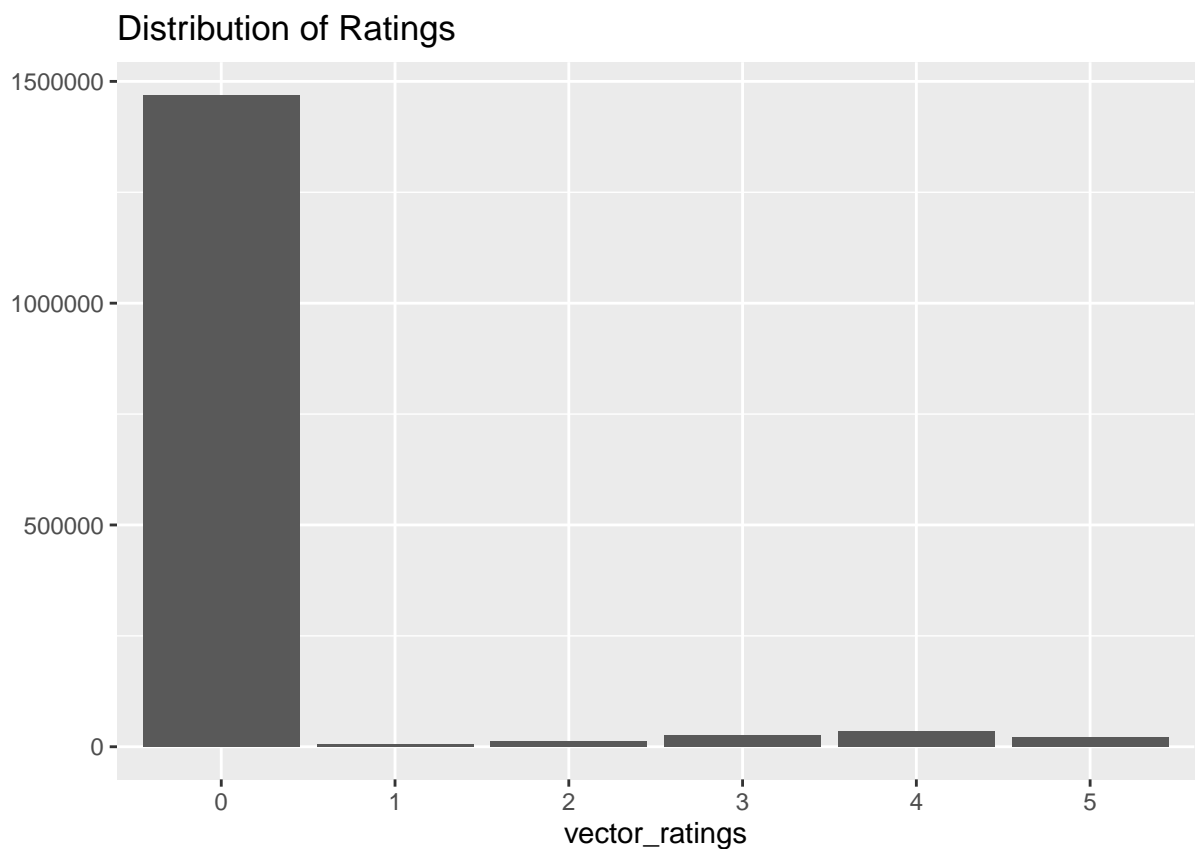
Table 1: Item Similarity Table

vector_ratings	Freq
0	1469760
1	6059
2	11307
3	27002
4	33947
5	21077

```
vector_ratings <- as.vector(matrix@data)
table_ratings <- table(vector_ratings)
table_ratings %>% kable(caption = "Item Similarity Table") %>% kable_styling("striped", full_width = TRUE)
```

We can look at the distribution of ratings.

```
vector_ratings <- factor(vector_ratings)
qplot(vector_ratings) + ggtitle("Distribution of Ratings")
```



### Select Relevant Data


We can subset this data to select users who have rated at least 95 movies, and movies that have been watched at least 50 times. The dimension of the matrix changes from 943 x 1664 to 369 x 691

```
movie_ratings <- matrix[rowCounts(matrix) > 95, colCounts(matrix) > 50]
dim(movie_ratings)
```

```
## [1] 369 591
```

Top 1% of Users

```
min_movies <- quantile(rowCounts(matrix, na.rm = TRUE), 0.99)
min_users <- quantile(colCounts(matrix, na.rm = TRUE), 0.99)
image(matrix[rowCounts(matrix) > min_movies, colCounts(matrix) > min_users], main = "Heatmap of the top
```



Distribution of Average Rating by User

```
average_rating <- rowMeans(movie_ratings)
qplot(average_rating) + stat_bin(binwidth = 0.2) + ggtitle("Distribution of Average Rating by User")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Distribution of Average Rating by User




Table 2: User 1 Recommendations: Item-Item

x
Six Degrees of Separation (1993)
While You Were Sleeping (1995)
Man Who Knew Too Little, The (1997)
Othello (1995)
Space Jam (1996)

## Split into Training & Test datasets using an 80/20 Ratio

```
set.seed(123)
train_sample <- sample(x = c(TRUE, FALSE), size = nrow(movie_ratings), replace = TRUE, prob = c(0.8, 0.2))
movie_train <- movie_ratings[train_sample,]
movie_test <- movie_ratings[!train_sample,]
```

## Item-Item Collaborative Filtering

This method identifies which items were similar in terms of other's items closely related.

```
item_model <- Recommender(movie_train, method = "IBCF")
item_model
```

### Create the Item-Item Collaborative Filter Recommender System with the training set

```
## Recommender of type 'IBCF' for 'realRatingMatrix'
## learned using 300 users.
```

```
item_predict <- predict(item_model, movie_test, n = 5)
```

**Predict the Item-Item Collaborative Filter Recommender System on the test set** Let's produce a list of the top 5 movie recommendations based on the Item-Item Collaborative Filter Recommender System.

```
item_reco_user_1 <- item_predict@items[[1]]
item_movie_user_1 <- item_predict@itemLabels[item_reco_user_1]
item_movie_user_1 %>% kable(caption = "User 1 Recommendations: Item-Item") %>% kable_styling("striped",
```

## User-User Collaborative Filtering

This method identifies similar user's to show close relations between the different users.



Table 3: User 1 Recommendations: User-User

x
Titanic (1997)
Good Will Hunting (1997)
Much Ado About Nothing (1993)
This Is Spinal Tap (1984)
Apt Pupil (1998)

```
user_model <- Recommender(movie_train, method = "UBCF")
user_model
```

```
## Recommender of type 'UBCF' for 'realRatingMatrix'
## learned using 300 users.
```

```
user_predict <- predict(user_model, movie_test, n = 5)
```

**Predict the User-User Collaborative Filter Recommender System on the test set** Let's produce a list of the top 5 movie recommendations based on the User-User Collaborative Filter Recommender System.

```
user_reco_user_1 <- user_predict@items[[1]]
user_movie_user_1 <- user_predict@itemLabels[user_reco_user_1]
user_movie_user_1 %>% kable(caption = "User 1 Recommendations: User-User") %>% kable_styling("striped",
```

## Summary

While both recommender systems recommended different movie's to the first user of the test set, each recommender system serves a different purpose. Both method's have their respective positives and negatives, but a hybrid approach might help create a better recommender system. The size of the data does affect performance, and the larger item-item collaborative system will produce less errors than the user-user collaborative system.

**Source Code** [GitHub Repository](#)